

Lecture notes: Week 7

This week, we began our explorations of recursion. We discussed:

- Principles of Recursion
- Recursively structured data
- Divide and conquer recursion
- Using GDB to explore recursion
- Merge sort

We also discussed mathematical induction, as a way of understanding why recursion works. You won't be tested on that, so

Recursion

In a superficial sense, recursion is what happens whenever a function f calls itself, as in the "standard" factorial example:

```
int factorial(int n) {  
    if (n == 0) return 1;  
    else return n * factorial(n - 1);  
}
```

It is better to think of recursion as a technique for solving problems. Many problems can be thought of using the following pattern:

1. Break the problem up into "smaller" version(s) of the same problem.
2. Solve the smaller problem(s) by calling the same function (we call this the *inductive* or *recursive* case)
3. Use the solutions to the smaller problem(s) to solve the original problem.

This seems like a process that doesn't end: to solve a big problem, we break it up into smaller versions of the problem -- but then we have to solve the smaller problem, which isn't any different! The key is that you can repeat this process, solving the smaller problems in the same way. At each step, you get smaller and smaller problems. Eventually, the problem is small enough that getting the answer is trivial. We call this the *base* case.

We can see this in the factorial example: rather than computing factorial of n , we realize that $n!$ is just $n * (n-1)!$ (Step 1: break the problem up into a smaller version of the same problem) -- so we can call $\text{factorial}(n - 1)$ (Step 2: solve the smaller problem by calling the same function). We can then multiply this by n to find $\text{factorial}(n)$ (Step 3: use the solution of the smaller problem to solve the original problem). We also see that the base case is simple: we already know what $0!$ is, so there is no need to "break it up" into a smaller problem -- we can just return 1.

(Note: you could also write factorial with a loop, and the loop version would probably be faster, so you might wonder why we need recursion. In some cases, recursion may be just an easier way to think of a problem than writing a loop. Some times recursion might even be much slower! But some times, recursion is the only way to effectively solve a problem.)

One way to think about how to correctly write a recursive function is to think *inductively*: We can assume that the recursive function already works, but only if the function is called on a smaller problem than what we're solving. We can then write the recursive function assuming that it already works. The only thing we have to make sure we do is write correct base cases -- we need to make sure that for the smallest versions of the function, we compute the correct answer.

We can see this strategy at work in the factorial code: when figuring out what `factorial(n)` should be, we can assume that `factorial(n-1)` already works, giving us $(n-1)!$ (even though we don't have a working factorial function yet). If we have $(n-1)!$, computing $n!$ is easy: just multiply by n . All that's left is to make sure that we have a working base case: that we just provide the answer for the "smallest" argument we will pass to factorial, in this case 0.

Recursively structured data

But programs like factorial are not a particularly interesting use of recursion. As mentioned above, factorial can just as easily be written with a loop. And other classic examples of recursion, like `fibonacci`, are even worse: a recursively-written fibonacci takes an exponential amount of time to run!

Where recursion really comes into its own is when we are thinking about *recursively structured data*. This is data where one of the fields (attributes) of the data has the same type as the data itself. We're used to seeing structures that define new data types. A Student data type (like in Homeworks 5 and 6) might look like this:

```
struct Student {
    int id;
    char firstname[80];
    char lastname[80];
};
```

Here, the data is not recursive: the fields of the data are of different types than the Student itself. But what if we want to create a program that deals with family trees? Well, one way of representing that data is to think about what the struct for a Person should be. They might have the same data as a Student (first name, last name, some identifier), but they also have two parents (please excuse the assumption that everyone has exactly two parents):

```
struct Person {
    int id;
    char firstname[80];
    char lastname[80];
    ____ parent1;
    ____ parent2;
};
```

Note that we're deliberately not showing the exact syntax for the data types of `parent1` and `parent2`. We'll discuss that in more detail when we get to our module on dynamic data structures. For now, it'll suffice to think that `parent1` and `parent2` are *also* `Person` structures: they have ids, first names, last names, and parents of their own!

For those of you who don't want to wait for the answer: the data type of `parent1` and `parent2` is `struct Person *`. A pointer to another `Person` structure! Note that these *have* to be pointers. But if we follow a `Person`'s `parent1` pointer, we'll get to another `Person` structure.

This is *recursively structured data*: Each `Person` has parent's who are also `Persons`, who have parents who are also `Persons`, etc. There may be a special `Person` structure we use for people who have unknown parent.

What if we want to solve a problem using this recursively structured data? For example, asking whether a person with id `12345` is my ancestor?

Well, we can think recursively! What makes someone my ancestor? Well, they're either one of my parents... or they're one of my parents' ancestors! So I can use this fact to write a recursive function that tries to determine whether `Person query` is `Person p`'s ancestor?

(Note: this is pseudocode, not real C):

```
bool isAncestor(Person p, Person query) {
    if (p's parent1 is not unknown) {
        // check if query is parent1
        if p's parent1 is query, return true
        // check if query is p's parent1's ancestor
        else if (isAncestor(p's parent1, query)) return true;
    }
    if (p's parent2 is not unknown) {
        // check if query is parent2
        if p's parent2 is query, return true;
        // check if query is p's parent2's ancestor
        else if (isAncestor(p's parent2, query)) return true;
    }
    // query is not p's parent, nor p's parents' ancestor
    // so is not p's ancestor!
    return false;
}
```

Note that this pseudocode uses the pattern we described above: we broke the problem down into smaller problems, each of which is either easy to solve (is query one of p's parents?) or is a smaller version of the problem we're trying to solve (is query an ancestor of p's parent1? Is query an ancestor of p's parent2?) We can then *assume that isAncestor works* to just call it again to solve the smaller problems: in this case, we call it twice, except to find ancestors of p's parents. We then construct the solution we want out of the solutions to the smaller problems.

Note that the base case—the smallest version of the problem—is handled by the fact that if a Person does not have parents, we'll fall through both if statements and just return false.

Divide-and-Conquer recursion

A very common pattern for recursive problems is divide-and-conquer recursion: to solve a problem on n pieces of data (e.g., an array of length n), we break the input up into two pieces, each with $n/2$ pieces of data (e.g., two arrays, each with half the elements), call the recursive function on these smaller pieces, then write some code to combine the results from those two functions into the final answer. The base case for this style of function is what to do when you have only 1 element.

Note that arguably, our `isAncestor` problem from above is an instance of this: we broke the family tree of Person `p` into two pieces: the family tree of `parent1`, and the family tree of `parent2`

Consider a toy example where we want to sum up all the values in an input array with n elements. Here, if we divide the array in two and sum those two sub-arrays, we can add the results to get the sum of the whole array. The base case is that the sum of an array with just one element is the value of that element:

```
int sum(int * arr, int nels) {
    if (nels == 1) return arr[0];

    int sum1 = sum(arr, nels/2);
    int sum2 = sum(&arr[nels/2], (nels + 1)/2);

    return sum1 + sum2
}
```

(The $(nels + 1)/2$ stuff is just a fancy way of dealing with arrays that have an odd number of elements, where `sum2` works over a slightly larger array than `sum1`. In integer division, $nels/2$ is like computing $\text{floor}(n/2)$, and $(nels + 1)/2$ is like computing $\text{ceiling}(n/2)$. More generally, to compute $\text{ceiling}(a/b)$ you can do integer division: $(a + b - 1)/b$.)

Note that this function makes use of the fact that arrays and pointers are treated similarly. When we want to sort the second half of the array, we can just pass the address of the beginning of the second half of the array (`&arr[nels/2]`) and the `sum` function will just treat it as an array that needs sorting—it doesn't care that it's the middle part of an existing array.

Exploring recursion with GDB

So what happens when we call a recursive function like `factorial`? Remember that each time we call a function, we push its *frame* on the stack, which maintains all of its local variables and arguments. So every time `factorial` calls itself, it pushes a new “copy” of `factorial` onto the stack, with its own copy of the argument `n`. This is the key to recursion: each call to the function has its

own local variables, which lets us multiply everything together as the functions return up the stack.

Let's see this in action using gdb. Let's say we compile this code:

```
#include <stdio.h>

int foo(int n) {
    int retval = n;

    if (n == 0) return 1;

    retval = retval * foo(n - 1);

    return retval;
}

void main() {
    int x = foo(5);
    printf("foo(5) = %d\n", x);
}
```

and then run it through gdb:

```
> gdb a.out
(gdb) r
Starting program: /home/dynamo/b/milind/264/recgdbtest/a.out
foo(5) = 120
```

We see the result of the code. Let's set a breakpoint in foo so we can investigate what is going on. But we will be clever: rather than just setting a breakpoint in foo ("b foo"), we'll set a breakpoint in foo that will *only* trigger when n is 0:

```
(gdb) b foo if (n == 0)
Breakpoint 1 at 0x4004cf: file rectest.c, line 5.
```

Now when we run the program, execution will stop when foo is called with argument 0:

```
(gdb) r
Starting program: /home/dynamo/b/milind/264/recgdbtest/a.out

Breakpoint 1, foo (n=0) at rectest.c:5
5      int retval = n;
```

If we step forward one step, we can see the value of retval:

```
(gdb) n
7      if (n == 0) return 1;
(gdb) p retval
$1 = 0
```

So what does the program stack look like? gdb can show us the *calling context* or *call stack* of the program at this point: the chain of functions that had to be called to get us to this point. To do this, we use the command “backtrace” (which can also be written “bt” or “where”):

```
(gdb) bt
#0  foo (n=0) at rectest.c:7
#1  0x00000000004004ef in foo (n=1) at rectest.c:9
#2  0x00000000004004ef in foo (n=2) at rectest.c:9
#3  0x00000000004004ef in foo (n=3) at rectest.c:9
#4  0x00000000004004ef in foo (n=4) at rectest.c:9
#5  0x00000000004004ef in foo (n=5) at rectest.c:9
#6  0x000000000040050f in main () at rectest.c:15
```

We see here the full “history” of the execution up to this point: main called foo(5) at line 15, foo then called foo(4) at line 9, foo then called foo(3) at line 9, and so on. Each one of those lines represents a *stack frame* sitting on the stack, each with its own local variables.

We can even jump to one of those stack frames to see what’s going on:

```
(gdb) f 2
#2  0x00000000004004ef in foo (n=2) at rectest.c:9
9      retval = retval * foo(n - 1);
```

Which jumps to the stack frame numbered 2 (frame 0 is always the frame of the currently running function, frame 1 is the frame of the function that called the current function, etc.)

Now if we inspect the value of retval, we’ll see that it has the value from a different version of foo:

```
(gdb) p retval
$2 = 2
```

And if we jump back to frame 0, we can see the “current” version of foo’s value of retval again:

```
(gdb) f 0
#0  foo (n=0) at rectest.c:7
7      if (n == 0) return 1;
(gdb) p retval
$3 = 0
```

Merge Sort

One of the classic examples of divide-and-conquer recursion is merge sort, which you have to write for HW6. The basic idea of merge sort is this. We divide the problem up as follows:

1. If the array only has 0 or 1 elements, it is already sorted (this is the base case)
2. Divide the array into two pieces (like in sum)
3. Sort the two pieces using merge sort (This is the recursive case. We can assume that this works properly!)
4. Now that we have two sorted half arrays, we will *merge* them together into a single sorted array—which is now the original array that was sorted!

So you can see that this is an instance of our recursion design pattern: we found solutions to the simplest versions of the problem (arrays with 0 or 1 element). Figured out how to break the problem up into smaller versions of the same problem (sorting half-arrays). Then combined the solutions of the smaller problem to solve the larger problem.

The magic of merge sort lies in how the *merge* is written. This is what you'll have to do for HW 6, so we won't give detailed examples here. But here's a useful way to think about it:

Suppose you have a shuffled deck of cards that you split in half, and then you sort the two halves of the deck. So now you have two half-decks of cards, each of which is sorted. How can you merge them into a single deck, which we'll call the *output* deck?

What you can do is look at the top card of each deck, and put the smaller of the two cards onto the output deck. Note that the output deck now definitely has the smallest card on it. And now we can repeat the process! Look at the top card of each deck and put the smaller card on the output deck. If we keep doing this until both decks are empty, we will have interleaved the two decks together onto the output deck.

In your actual merge, you don't have decks of cards. You have three arrays: the two half arrays, and the output array. Think about maintaining three pointers: a pointer to where you currently are in the output array (how many cards have you put on the output array?) And pointers to where you currently are in the two half-arrays (how many cards have you already removed from each half-array?) Just walk those pointers through the arrays until you have moved all the data from the half arrays to the output array.